

EX.NO:

DATE:

A MULTILAYER PERCEPTRON WITH A HYPERPARAMETER TUNING

AIM:

To implement a Multilayer Perceptron (MLP) model and perform hyperparameter tuning to improve its performance on a given dataset.

ALGORITHM:

STEP1 : Import all necessary libraries including pandas, NumPy, scikit-learn modules, and TensorFlow Keras components for building and evaluating the neural network model.

STEP2 : Load the dataset using pandas and read it into a DataFrame from the specified file path.

STEP3 : Create a binary target column named 'pass' by checking if the GPA is greater than or equal to 2.0.

STEP4 : Engineer a new feature called study_attendance by dividing weekly study time to reflect effort adjusted for attendance.

STEP5 : Remove irrelevant or redundant columns such as StudentID, GPA, and GradeClass to avoid data leakage and unnecessary noise.

STEP6 : Apply one-hot encoding to convert categorical variables into numeric format using dummy variables while dropping the first category to avoid multicollinearity.

STEP7 : Separate the input features (X) and the target variable (y) for modeling.

STEP8 : Normalize the feature set using StandardScaler to ensure that all features contribute equally to the training process.

STEP9 : Split the data into training and testing sets using an 80-20 ratio and define class weights to address any class imbalance.

STEP10 : Build and compile a Sequential MLP model with multiple hidden layers and dropout, train it with early stopping, and evaluate the model's accuracy and classification report on the test set.

SOURCE CODE:

```
# Import libraries

import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler


# 2. Load dataset

df = pd.read_csv('/content/drive/MyDrive/Student_performance_data_.csv')


# 3. Create binary target: pass (1 if GPA >= 2.0)

df['pass'] = (df['GPA'] >= 2.0).astype(int)


# 4. Feature engineering

df['study_attendance'] = df['StudyTimeWeekly'] / (1 + df['Absences'])


# 5. Drop unnecessary columns

df = df.drop(['StudentID', 'GPA', 'GradeClass'], axis=1)


# 6. One-hot encoding for categorical features

df = pd.get_dummies(df, drop_first=True)


# 7. Split features and target

X = df.drop('pass', axis=1)
y = df['pass']


# 8. Scale features

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)


# 9. Train-test split
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

```
# 1. Import Libraries
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Dense, Dropout
```

```
from tensorflow.keras.optimizers import Adam
```

```
from tensorflow.keras.callbacks import EarlyStopping
```

```
# 2. Load Dataset
```

```
df = pd.read_csv('/content/drive/MyDrive/Student_performance_data_.csv')
```

```
# 3. Create Binary Target: Pass if GPA >= 2.0
```

```
df['pass'] = (df['GPA'] >= 2.0).astype(int)
```

```
# 4. Feature Engineering
```

```
df['study_attendance'] = df['StudyTimeWeekly'] / (1 + df['Absences'])
```

```
# 5. Drop Unnecessary Columns
```

```
df = df.drop(['StudentID', 'GPA', 'GradeClass'], axis=1)
```

```
# 6. One-Hot Encoding for Categorical Features
```

```
df = pd.get_dummies(df, drop_first=True)
```

```
# 7. Split Features and Target
```

```
X = df.drop('pass', axis=1)
```

```
y = df['pass']
```

8. Feature Scaling

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

9. Train-Test Split

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

10. Build MLP Model with 3 Hidden Layers

```
model = Sequential([
    Dense(128, activation='relu', input_shape=(X.shape[1],)),
    Dropout(0.3),
    Dense(64, activation='relu'),
    Dropout(0.2),
    Dense(32, activation='relu'), # Additional Hidden Layer
    Dense(1, activation='sigmoid') # Binary output
])
```

Compile Model

```
model.compile(optimizer=Adam(learning_rate=0.0005),
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

11. EarlyStopping to Prevent Overfitting

```
early_stop = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
```

12. Train the Model

```
history = model.fit(X_train, y_train,
                    validation_split=0.2,
                    epochs=100,
                    batch_size=16,
```

```
callbacks=[early_stop],

verbose=1)
```

13. Evaluate Final Test Accuracy

```
test_loss, test_acc = model.evaluate(X_test, y_test, verbose=0)

print(f"\n👍 Final Test Accuracy: {test_acc * 100:.2f}%")
```

14. Classification Report

```
y_pred = (model.predict(X_test) > 0.5).astype("int32")

print("\nClassification Report:")

print(classification_report(y_test, y_pred))
```

OUTPUT

```
Epoch 1/100
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
96/96 ━━━━━━━━━━━ 2s 3ms/step - accuracy: 0.6549 - loss: 0.6490 - val_accuracy: 0.9080 - val_loss: 0.4043
Epoch 2/100
96/96 ━━━━━━━━━━━ 0s 3ms/step - accuracy: 0.8873 - loss: 0.3683 - val_accuracy: 0.9347 - val_loss: 0.1701
Epoch 3/100
96/96 ━━━━━━━━━━━ 1s 3ms/step - accuracy: 0.9187 - loss: 0.1907 - val_accuracy: 0.9470 - val_loss: 0.1383
Epoch 4/100
96/96 ━━━━━━━━━━━ 0s 3ms/step - accuracy: 0.9347 - loss: 0.1622 - val_accuracy: 0.9556 - val_loss: 0.1275
Epoch 5/100
96/96 ━━━━━━━━━━━ 1s 3ms/step - accuracy: 0.9233 - loss: 0.1704 - val_accuracy: 0.9608 - val_loss: 0.1206
Epoch 6/100
96/96 ━━━━━━━━━━━ 1s 3ms/step - accuracy: 0.9361 - loss: 0.1468 - val_accuracy: 0.9594 - val_loss: 0.1242
Epoch 7/100
96/96 ━━━━━━━━━━━ 0s 3ms/step - accuracy: 0.9453 - loss: 0.1442 - val_accuracy: 0.9582 - val_loss: 0.1214
Epoch 8/100
96/96 ━━━━━━━━━━━ 1s 3ms/step - accuracy: 0.9409 - loss: 0.1366 - val_accuracy: 0.9470 - val_loss: 0.1234
Epoch 9/100
96/96 ━━━━━━━━━━━ 0s 3ms/step - accuracy: 0.9434 - loss: 0.1367 - val_accuracy: 0.9504 - val_loss: 0.1204
Epoch 10/100
96/96 ━━━━━━━━━━━ 0s 3ms/step - accuracy: 0.9431 - loss: 0.1266 - val_accuracy: 0.9530 - val_loss: 0.1214
Epoch 11/100
96/96 ━━━━━━━━━━━ 1s 5ms/step - accuracy: 0.9430 - loss: 0.1334 - val_accuracy: 0.9556 - val_loss: 0.1190
Epoch 12/100
96/96 ━━━━━━━━━━━ 1s 5ms/step - accuracy: 0.9419 - loss: 0.1309 - val_accuracy: 0.9582 - val_loss: 0.1239
Epoch 13/100
96/96 ━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.9559 - loss: 0.1181 - val_accuracy: 0.9470 - val_loss: 0.1235
Epoch 14/100
96/96 ━━━━━━━━━━━ 1s 5ms/step - accuracy: 0.9515 - loss: 0.1090 - val_accuracy: 0.9470 - val_loss: 0.1259
Epoch 15/100
96/96 ━━━━━━━━━━━ 1s 5ms/step - accuracy: 0.9356 - loss: 0.1274 - val_accuracy: 0.9470 - val_loss: 0.1242
Epoch 16/100
96/96 ━━━━━━━━━━━ 1s 5ms/step - accuracy: 0.9423 - loss: 0.1309 - val_accuracy: 0.9504 - val_loss: 0.1265

👍 Final Test Accuracy: 93.11%
15/15 ━━━━━━━━━━━ 0s 5ms/step

Classification Report:
      precision    recall  f1-score   support

0               0.93       0.94       0.93         249
1               0.93       0.93       0.93         230

accuracy               0.93
macro avg              0.93       0.93       0.93         479
weighted avg           0.93       0.93       0.93         479
```

1. Setup

```
import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.metrics import classification_report, accuracy_score
```

```
import tensorflow as tf

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Dense, Dropout

from tensorflow.keras.callbacks import EarlyStopping


# 2. Load Data

df = pd.read_csv('/content/drive/MyDrive/Student_performance_data_.csv')


# 3. Create binary target: pass if GPA >= 2.0

df['pass'] = (df['GPA'] >= 2.0).astype(int)


# 4. Feature Engineering

df['study_attendance'] = df['StudyTimeWeekly'] / (1 + df['Absences'])


# 5. Drop Unused Columns

df.drop(['StudentID', 'GPA', 'GradeClass'], axis=1, inplace=True)


# 6. One-hot Encoding

df = pd.get_dummies(df, drop_first=True)


# 7. Define Features & Target

X = df.drop('pass', axis=1)

y = df['pass']


# 8. Normalize Features

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)


# 9. Train/Test Split
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

```
# 10. Handle Imbalance (adjust if needed based on your class distribution)
```

```
class_weights = {0: 1.2, 1: 0.8}
```

```
# 11. Build Model
```

```
model = Sequential([  
    Dense(256, activation='relu', input_shape=(X.shape[1],)),  
    Dense(128, activation='relu'),  
    Dense(64, activation='relu'),  
    Dropout(0.2),  
    Dense(1, activation='sigmoid')  
])
```

```
# 12. Compile Model
```

```
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.0005),  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

```
# 13. EarlyStopping
```

```
early_stop = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)
```

```
# 14. Train Model
```

```
history = model.fit(X_train, y_train,  
                    epochs=100,  
                    batch_size=32,  
                    validation_split=0.2,  
                    class_weight=class_weights,  
                    callbacks=[early_stop],  
                    verbose=1)
```

15. Evaluate Accuracy

```
loss, acc = model.evaluate(X_test, y_test, verbose=0)

print(f"\n✔ Improved Test Accuracy: {acc * 100:.2f}%")
```

16. Classification Report

```
y_pred = (model.predict(X_test) > 0.5).astype(int)

print("\nImproved Classification Report:")

print(classification_report(y_test, y_pred))
```

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the first layer in the model.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/100
48/48 ----- 6s 8ms/step - accuracy: 0.6171 - loss: 0.5988 - val_accuracy: 0.8938 - val_loss: 0.2793
Epoch 2/100
48/48 ----- 0s 4ms/step - accuracy: 0.9232 - loss: 0.2057 - val_accuracy: 0.9321 - val_loss: 0.1745
48/48 ----- 0s 4ms/step - accuracy: 0.9399 - loss: 0.1275 - val_accuracy: 0.9478 - val_loss: 0.1337
Epoch 4/100
48/48 ----- 0s 4ms/step - accuracy: 0.9482 - loss: 0.1157 - val_accuracy: 0.9399 - val_loss: 0.1543
Epoch 5/100
48/48 ----- 0s 5ms/step - accuracy: 0.9562 - loss: 0.1091 - val_accuracy: 0.9426 - val_loss: 0.1386
Epoch 6/100
48/48 ----- 0s 4ms/step - accuracy: 0.9553 - loss: 0.0976 - val_accuracy: 0.9478 - val_loss: 0.1352
Epoch 7/100
48/48 ----- 0s 4ms/step - accuracy: 0.9580 - loss: 0.0971 - val_accuracy: 0.9478 - val_loss: 0.1368
Epoch 8/100
48/48 ----- 0s 4ms/step - accuracy: 0.9717 - loss: 0.0834 - val_accuracy: 0.9478 - val_loss: 0.1387
Epoch 9/100
48/48 ----- 0s 4ms/step - accuracy: 0.9653 - loss: 0.0845 - val_accuracy: 0.9452 - val_loss: 0.1434
Epoch 10/100
48/48 ----- 0s 4ms/step - accuracy: 0.9700 - loss: 0.0721 - val_accuracy: 0.9556 - val_loss: 0.1383
Epoch 11/100
48/48 ----- 0s 5ms/step - accuracy: 0.9655 - loss: 0.0844 - val_accuracy: 0.9321 - val_loss: 0.1658
Epoch 12/100
48/48 ----- 0s 6ms/step - accuracy: 0.9676 - loss: 0.0797 - val_accuracy: 0.9399 - val_loss: 0.1634
Epoch 13/100
48/48 ----- 1s 6ms/step - accuracy: 0.9698 - loss: 0.0738 - val_accuracy: 0.9347 - val_loss: 0.1676

✔ Improved Test Accuracy: 93.11%
13/15 ----- 0s 7ms/step

Improved Classification Report:
precision    recall  f1-score   support

0           0.92       0.96       0.94       249
1           0.95       0.90       0.93       230

accuracy          0.93          0.93          0.93       479
macro avg         0.93          0.93          0.93       479
weighted avg      0.93          0.93          0.93       479

```

```
from sklearn.metrics import confusion_matrix
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
```

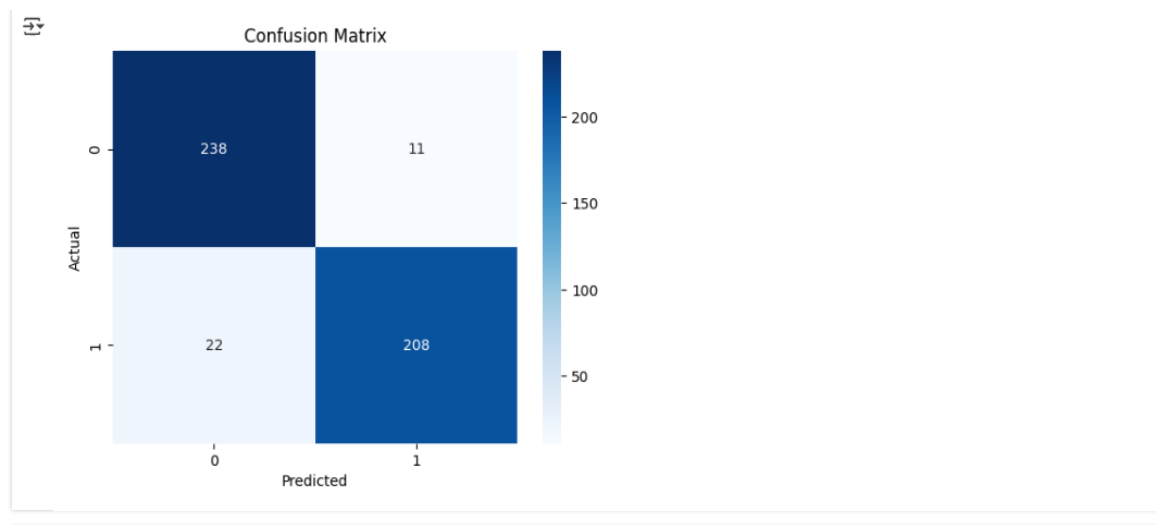
```
plt.xlabel("Predicted")
```

```
plt.ylabel("Actual")
```

```
plt.title("Confusion Matrix")
```

```
plt.show()
```


OUTPUT



COE(20)	
RECORD(20)	
VIVA(10)	
TOTAL(50)	

RESULT:

The Multilayer Perceptron model achieved high accuracy in classifying student performance based on the provided features.